On Performance Stability in LSM-based Storage Systems

Chen Luo
University of California, Irvine
cluo8@uci.edu

Michael J. Carey
University of California, Irvine
mjcarey@ics.uci.edu

ABSTRACT

The Log-Structured Merge-Tree (LSM-tree) has been widely adopted for use in modern NoSQL systems for its superior write performance. Despite the popularity of LSM-trees, they have been criticized for suffering from write stalls and large performance variances due to the inherent mismatch between their fast in-memory writes and slow background I/O operations. In this paper, we use a simple yet effective two-phase experimental approach to evaluate write stalls for various LSM-tree designs. We further identify and explore the design choices of LSM merge schedulers to minimize write stalls given an I/O bandwidth budget. We have conducted extensive experiments in the context of the Apache AsterixDB system and we present the results here.

1. INTRODUCTION

In recent years, the Log-Structured Merge-Tree (LSM-tree) [44, 40] has been widely used in modern key-value stores and NoSQL systems [2, 4, 5, 7, 14]. Different from traditional index structures, such as B^+-trees, which apply updates in-place, an LSM-tree always buffers writes into memory. When memory is full, writes are flushed to disk and subsequently merged using sequential I/Os. To improve efficiency and minimize blocking, flushes and merges are often performed asynchronously in the background.

Despite their popularity, LSM-trees have been criticized for suffering from write stalls and large performance variances [3, 50, 56]. To illustrate this problem, we conducted a micro-experiment on RocksDB [7], a state-of-the-art LSM-based key-value store, to evaluate its write throughput on SSDs using the YCSB benchmark [24]. The instantaneous write throughput over time is depicted in Figure 1. As one can see, the write throughput of RocksDB periodically slows down after the first 300 seconds, which is when the system

Figure 1: **Instantaneous write throughput of RocksDB: writes are periodically stalled to wait for lagging merges**

has to wait for background merges to catch up. Write stalls can significantly impact percentile write latencies and must be minimized to improve the end-user experience or to meet strict service-level agreements [35].

In this paper, we study the impact of write stalls and how to minimize write stalls for various LSM-tree designs. It should first be noted that some write stalls are inevitable. Due to the inherent mismatch between fast in-memory writes and slower background I/O operations, in-memory writes must be slowed down or stopped if background flushes or merges cannot catch up. Without such a flow control mechanism, the system will eventually run out of memory (due to slow flushes) or disk space (due to slow merges). Thus, it is not a surprise that an LSM-tree can exhibit large write stalls if one measures its maximum write throughput by writing as quickly data as possible, such as we did in Figure 1.

This inevitability of write stalls does not necessarily limit the applicability of LSM-trees since in practice writes do not arrive as quickly as possible, but rather are controlled by the expected data arrival rate. The data arrival rate directly impacts the write stall behavior and resulting write latencies of an LSM-tree. If the data arrival rate is relatively low, then write stalls are unlikely to happen. However, it is also desirable to maximize the supported data arrival rate so that the system's resources can be fully utilized. Moreover, the expected data arrival rate is subject to an important constraint - it must be smaller than the processing capacity of the target LSM-tree. Otherwise, the LSM-tree will never be able to process writes as they arrive, causing infinite write latencies. Thus, to evaluate the write stalls of an LSM-tree, the first step is to choose a proper data arrival rate.

As the first contribution, we propose a simple yet effective approach to evaluate the write stalls of various LSM-tree designs by answering the following question: If we set the

data arrival rate close to (e.g., 95% of) the maximum write throughput of an LSM-tree, will that cause write stalls? In other words, can a given LSM-tree design provide both a high write throughput and a low write latency? Briefly, the proposed approach consists of two phases: a *testing* phase and a *running* phase. During the testing phase, we experimentally measure the maximum write throughput of an LSM-tree by simply writing as much data as possible. During the running phase, we then set the data arrival rate close to the measured maximum write throughput as the limiting data arrival rate to evaluate its write stall behavior based on write latencies. If write stalls happen, the measured write throughput is not *sustainable* since it cannot be used in the long-term due to the large latencies. However, if write stalls do not happen, then write stalls are no longer a problem since the given LSM-tree can provide a high write throughput with small performance variance.

Although this approach seems to be straightforward at first glance, there exist two challenges that must be addressed. First, how can we accurately measure the maximum sustainable write rate of an LSM-tree experimentally? Second, how can we best schedule LSM I/O operations so as to minimize write stalls at runtime? In the remainder of this paper, we will see that the merge scheduler of an LSM-tree can have a large impact on write stalls. As the paper's second contribution, we identify and explore the design choices for LSM merge schedulers and present a new merge scheduler to address these two challenges.

As our final contribution, we have implemented the proposed techniques and various LSM-tree designs inside Apache AsterixDB [14]. This enabled us to carry out extensive experiments to evaluate the write stalls of LSM-trees and the effectiveness of the proposed techniques using our two-phase evaluation approach. We argue that with proper tuning and configuration, LSM-trees can achieve both a high write throughput and small performance variance.

The remainder of this paper is organized as follows: Section 2 provides background on LSM-trees and discusses related work. Section 3 describes the general experimental setup used throughout this paper. Section 4 identifies the scheduling choices for LSM-trees and evaluates bLSM's merge scheduler [50]. Sections 5 and 6 present our techniques for minimizing write stalls for various LSM-tree designs. Section 7 summarize the important lessons and insights from our evaluation. Finally, Section 8 concludes the paper. An extended version of this paper [41] further extends our evaluation to the size-tiered merge policy used in practical systems and LSM-based secondary indexes.

2. BACKGROUND

2.1 Log-Structured Merge Trees

The LSM-tree [44] is a persistent index structure optimized for write-intensive workloads. In an LSM-tree, writes are first buffered into a memory component. An insert or update simply adds a new entry with the same key, while a delete adds an anti-matter entry indicating that a key has been deleted. When the memory component is full, it is flushed to disk to form a new disk component, within which entries are ordered by keys. Once flushed, LSM disk components are immutable.

A query over an LSM-tree has to reconcile the entries with identical keys from multiple components, as entries


Figure 2: LSM-tree Merge Policies

from newer components override those from older components. A point lookup query simply searches all components from newest to oldest until the first match is found. A range query searches all components simultaneously using a priority queue to perform reconciliation. To speed up point lookups, a common optimization is to build Bloom filters [19] over the sets of keys stored in disk components. If a Bloom filter reports that a key does not exist, then that disk component can be excluded from searching. As disk components accumulate, query performance tends to degrade since more components must be examined. To counter this, smaller disk components are gradually merged into larger ones. This is implemented by scanning old disk components to create a new disk component with unique entries. The decision of what disk components to merge is made by a pre-defined *merge policy*, which is discussed below.

Merge Policy. Two types of LSM merge policies are commonly used in practice [40], both of which organize components into "levels". The leveling merge policy (Figure 2a) maintains one component per level, and a component at Level $i + 1$ will be T times larger than that of Level i. As a result, the component at Level i will be merged multiple times with the component from Level $i - 1$ until it fills up and is then merged into Level $i + 1$. In contrast, the tiering merge policy (Figure 2b) maintains multiple components per level. When a Level i becomes full with T components, these T components are merged together into a new component at Level $i + 1$. In both merge policies, T is called the *size ratio*, as it controls the maximum capacity of each level. We will refer to both of these merge policies as *full merges* since components are merged entirely.

In general, the leveling merge policy optimizes for query performance by minimizing the number of components but at the cost of write performance. This design also maximizes space efficiency, which measures the amount of space used for storing obsolete entries, by having most of the entries at the largest level. In contrast, the tiering merge policy is more write-optimized by reducing the merge frequency, but this leads to lower query performance and space utilization.

Partitioning. Partitioning is a commonly used optimization in modern LSM-based key-value stores that is often implemented together with the leveling merge policy, as pioneered by LevelDB [5]. In this optimization, a large LSM disk component is range-partitioned into multiple (often fixed-size) files. This bounds the processing time and the temporary space of each merge. An example of a partitioned LSM-tree with the leveling merge policy is shown in Figure 3, where each file is labeled with its key range. Note that partitioning starts from Level 1, as components in Level 0 are directly flushed from memory. To merge a file from Level i to Level $i + 1$, all of its overlapping files at

Figure 3: Partitioned LSM-tree with Leveling Merge Policy

Figure 4: bLSM's Spring-and-Gear Merge Scheduler

Level $i + 1$ are selected and these files are merged to form new files at Level $i + 1$. For example in Figure 3, the file labeled 0-50 at Level 1 will be merged with the files labeled 0-20 and 22-52 at Level 2, which produce new files labeled 0-15, 17-30, and 32-52 at Level 2. To select which file to merge next, LevelDB uses a round-robin policy.

Both full merges and partitioned merges are widely used in existing systems. For example, full merges are used in AsterixDB [1], Cassandra [2], HBase [4], ScyllaDB [8], Tarantool [9], and WiredTiger (MongoDB) [11]. Partitioned merges are used in LevelDB [5], RocksDB [7], and X-Engine [33].

Write Stalls in LSM-trees. Since in-memory writes are inherently faster than background I/O operations, writing to memory components sometimes must be stalled to ensure the stability of an LSM-tree, which, however, will negatively impact write latencies. This is often referred to as the *write stall* problem. If the incoming write speed is faster than the flush speed, writes will be stalled when all memory components are full. Similarly, if there are too many disk components, new writes should be stalled as well. In general, merges are the major source of write stalls since writes are flushed once but merged multiple times. Moreover, flush stalls can be avoided by giving higher I/O priority to flushes. In this paper, we thus focus on write stalls caused by merges.

2.2 Apache AsterixDB

Apache AsterixDB [1, 14, 22] is a parallel, semi-structured Big Data Management System (BDMS) that aims to manage large amounts of data efficiently. It supports a feed-based framework for efficient data ingestion [31, 54]. The records of a dataset in AsterixDB are hash-partitioned based on their primary keys across multiple nodes of a shared-nothing cluster; thus, a range query has to search all nodes. Each partition of a dataset uses a primary LSM-based B^+-tree index to store the data records, while local secondary indexes, including LSM-based B^+-trees, R-trees, and inverted indexes, can be built to expedite query processing. AsterixDB internally uses a variation of the tiering merge policy to manage disk components, similar to the one used in existing systems [4, 7]. Instead of organizing components into levels explicitly as in Figure 2b, AsterixDB's variation simply schedules merges based on the sizes of disk components. In this work, we do not focus on the LSM-tree implementation in AsterixDB but instead use AsterixDB as a common testbed to evaluate various LSM-tree designs.

2.3 Related Work

LSM-trees. Recently, a large number of improvements of the original LSM-tree [44] have been proposed. Luo and Carey [40] survey these improvements, ranging from improving write performance [18, 27, 28, 38, 43, 46, 55], optimizing

memory management [13, 17, 51, 57], supporting automatic tuning of LSM-trees [25, 26, 37], optimizing LSM-based secondary indexes [39, 45], to extending the applicability of LSM-trees [42, 48]. However, all of these efforts have largely ignored performance variances and write stalls.

Several LSM-tree implementations seek to bound the write processing latency to alleviate the negative impact of write stalls [5, 7, 36, 56]. bLSM [50] proposes a spring-and-gear merge scheduler to avoid write stalls. As shown in Figure 4, bLSM has one memory component, C_0, and two disk components, C_1 and C_2. The memory component C_0 is continuously flushed and merged with C_1. When C_1 becomes full, a new C_1 component is created while the old C_1, which now becomes C_1', will be merged with C_2. bLSM ensures that for each Level i, the progress of merging C_i' into C_{i+1} (denoted as "out_i") will be roughly identical to the progress of the formation of a new C_i (denoted as "in_i"). This eventually limits the write rate for the memory component (in_0) and avoids blocking writes. However, we will see later that simply bounding the maximum write processing latency alone is insufficient, because a large variance in the processing rate can still cause large queuing delays for subsequent writes.

Performance Stability. Performance stability has long been recognized as a critical performance metric. The TPC-C benchmark [10] measures not only absolute throughput, but also specifies the acceptable upper bounds for the percentile latencies. Huang et al. [35] applied VProfiler [34] to identify major sources of variance in database transactions. Various techniques have been proposed to optimize the variance of query processing [15, 16, 20, 23, 47, 53]. Cao et al. [21] found that variance is common in storage stacks and heavily depends on configurations and workloads. Dean and Barroso [29] discussed several engineering techniques to reduce performance variances at Google. Different from these efforts, in this work we focus on the performance variances of LSM-trees due to their inherent out-of-place update design.

3. EXPERIMENTAL METHODOLOGY

For ease of presentation, we will mix our techniques with a detailed performance analysis for each LSM-tree design. We now describe the general experimental setup and methodology for all experiments to follow.

3.1 Experimental Setup

All experiments were run on a single node with an 8-core Intel i7-7567U 3.5GHZ CPU, 16 GB of memory, a 500GB SSD, and a 1TB 7200 rpm hard disk. We used the SSD for LSM storage and configured the hard disk for transaction logging due to its sufficiently high sequential throughput. We allocated 10GB of memory for the AsterixDB instance. Within that allocation, the buffer cache size was set at 2GB. Each LSM memory component had a 128MB budget, and each LSM-tree had two memory components to minimize stalls during flushes. Each disk component had a Bloom

filter with a false positive rate setting of 1%. The data page size was set at 4KB to align with the SSD page size.

It is important to note that not all sources of performance variance can be eliminated [35]. For example, writing a key-value pair with a 1MB value inherently requires more work than writing one that only has 1KB. Moreover, short time periods with quickly occurring writes (workload bursts) will be much more likely to cause write stalls than a long period of slow writes, even though their long-term write rate may be the same. In this paper, we will focus on the *avoidable variance* [35] caused by the internal implementation of LSM-trees instead of variances in the workloads.

To evaluate the internal variances of LSM-trees, we adopt YCSB [24] as the basis for our experimental workload. Instead of using the pre-defined YCSB workloads, we designed our own workloads to better study the performance stability of LSM-trees. Each experiment first loads an LSM-tree with 100 million records, in random key order, where each record has size 1KB. It then runs for 2 hours to update the previously loaded LSM-tree. This ensures that the measured write throughput of an LSM-tree is stable over time. Unless otherwise noted, we used one writer thread for writing data to the LSM memory components. We evaluated two update workloads, where the updated keys follow either a uniform or Zipf distribution. The specific workload setups will be discussed in the subsequent sections.

We used two commonly used I/O optimizations when implementing LSM-trees, namely I/O throttling and periodic disk forces. In all experiments, we throttled the SSD write speed of all LSM flush and merge operations to 100MB/s. This was implemented by using a rate limiter to inject artificial sleeps into SSD writes. This mechanism bounds the negative impact of the SSD writes on query performance and allows us to more fairly compare the performance differences of various LSM merge schedulers. We further had each flush or merge operation force its SSD writes after each 16MB of data. This helps to limit the OS I/O queue length, reducing the negative impact of SSD writes on queries. We have verified that disabling this optimization would not impact the performance trends of writes; however, large forces at the end of each flush and merge operation, which are required for durability, can significantly interfere with queries.

3.2 Performance Metrics

To quantify the impact of write stalls, we will not only present the write throughput of LSM-trees but also their write latencies. However, there are different models for measuring write latencies. Throughout the paper, we will use *arrival rate* to denote the rate at which writes are submitted by clients, *processing rate* to denote the rate at which writes can be processed by an LSM-tree, and *write throughput* to denote the number of writes processed by an LSM-tree per unit of time. The difference between the write throughput and arrival/processing rates is discussed further below.

The bLSM paper [50], as well as most of the existing LSM research, used the experimental setup depicted in Figure 5a to write as much data as possible and measure the latency of each write. In this *closed system* setup [32], the processing rate essentially controls the arrival rate, which further equals the write throughput. Although this model is sufficient for measuring the maximum write throughput of LSM-trees, it is not suitable for characterizing their write latencies for several reasons. First, writing to memory is inherently faster


Figure 5: Models for Measuring Write Latency

than background I/Os, so an LSM-tree will always have to stall writes in order to wait for lagged flushes and merges. Moreover, under this model, a client cannot submit its next write until its current write is completed. Thus, when the LSM-tree is stalled, only a small number of ongoing writes will actually experience a large latency since the remaining writes have not been submitted yet[1].

In practice, a DBMS generally cannot control how quickly writes are submitted by external clients, nor will their writes always arrive as fast as possible. Instead, the arrival rate is usually independent from the processing rate, and when the system is not able to process writes as fast as they arrive, the newly arriving writes must be temporarily queued. In such an *open system* (Figure 5b), the measured write latency includes both the queuing latency and processing latency. Moreover, an important constraint is that the arrival rate must be smaller than the processing rate since otherwise the queue length will be unbounded. Thus, the (overall) write throughput is actually determined by the arrival rate.

A simple example will illustrate the important difference between these two models. Suppose that 5 clients are used to generate an intended arrival rate of 1000 writes/s and that the LSM-tree stalls for 1 second. Under the closed system model (Figure 5a), only 5 delayed writes will experience a write latency of 1s since the remaining (intended) 995 writes simply will not occur. However, under the open system model (Figure 5b), all 1000 writes will be queued and their average latency will be at least 0.5s.

To evaluate write latencies in an open system, one must first set the arrival rate properly since the write latency heavily depends on the arrival rate. It is also important to maximize the arrival rate to maximize the system's utilization. For these reasons, we propose a two-phase evaluation approach with a *testing* phase and a *running* phase. During the *testing* phase, we use the closed system model (Figure 5a) to measure the maximum write throughput of an LSM-tree, which is also its processing rate. When measuring the maximum write throughput, we excluded the initial 20-minute period (out of 2 hours) of the testing phase since the initially loaded LSM-tree has a relatively small number of disk components at first. During the *running* phase, we use the open system model (Figure 5b) to evaluate the *write latencies* under a constant arrival rate set at 95% of the measured maximum write throughput. Based on queuing theory [32], the queuing time approaches infinity when the utilization, which is the ratio between the arrival rate and the processing rate, approaches 100%. We thus empirically determine a high utilization load (95%) while leaving some room for the system to absorb variance. If the running phase then reports large write latencies, the maximum write

---

[1] The original release of the YCSB benchmark [24] mistakenly used this model; this was corrected later in 2015 [12].

throughput as determined in the testing phase is not sustainable; we must improve the implementation of the LSM-tree or reduce the expected arrival rate to reduce the latencies. In contrast, if the measured write latency is small, then the given LSM-tree can provide a high write throughput with a small performance variance.

4. LSM MERGE SCHEDULER

Different from a merge policy, which decides which components to merge, a *merge scheduler* is responsible for executing the merge operations created by the merge policy. In this section, we discuss the design choices for a merge scheduler and evaluate bLSM's spring-and-gear merge scheduler.

4.1 Scheduling Choices

The write cost of an LSM-tree, which is the number of I/Os per write, is determined by the LSM-tree design itself and the workload characteristics but not by how merges are executed [25]. Thus, a merge scheduler will have little impact on the overall write throughput of an LSM-tree as long as the allocated I/O bandwidth budget can be fully utilized. However, different scheduling choices can significantly impact the write stalls of an LSM-tree, and merge schedulers must be carefully designed to minimize write stalls. We have identified the following design choices for a merge scheduler.

Component Constraint: A merge scheduler usually specifies an upper-bound constraint on the total number of components allowed to accumulate before incoming writes to the LSM memory components should be stalled. We call this the *component constraint*. For example, bLSM [50] allows at most two disk components per level, while other systems like HBase [4] or Cassandra [2] specify the total number of disk components across all levels.

Interaction with Writes: There exist different strategies to enforce a given component constraint. One strategy is to simply stop processing writes once the component constraint is violated. Alternatively, the processing of writes can be degraded gracefully based on the merge pressure [50].

Degree of Concurrency: In general, an LSM-tree can often create multiple merge operations in the same time period. A merge scheduler should decide how these merge operations should be scheduled. Allowing concurrent merges will enable merges at multiple levels to proceed concurrently, but they will also compete for CPU and I/O resources, which can negatively impact query performance [13]. As two examples, bLSM [50] allows one merge operation per level, while LevelDB [5] uses just one single background thread to execute all merges one by one.

I/O Bandwidth Allocation: Given multiple concurrent merge operations, the merge scheduler should further decide how to allocate the available I/O bandwidth among these merge operations. A commonly used heuristic is to allocate I/O bandwidth "fairly" (evenly) to all active merge operations. Alternatively, bLSM [50] allocates I/O bandwidth based on the relative progress of the merge operations to ensure that merges at each level all make steady progress.

4.2 Evaluation of bLSM

Due to the implementation complexity of bLSM and its dependency on a particular storage system, Stasis [49], we chose to directly evaluate the released version of bLSM [6]. bLSM uses the leveling merge policy with two on-disk levels. We set its memory component size to 1GB and size ratio to 10 so that the experimental dataset with 100 million records can fit into the last level. We used 8 write threads to maximize the write throughput of bLSM.

Testing Phase. During the testing phase, we measured the maximum write throughput of bLSM by writing as much data as possible using both the uniform and Zipf update workloads. The instantaneous write throughput of bLSM under these two workloads is shown in Figure 6a. For readability, the write throughput is averaged over 30-second windows. (Unless otherwise noted, the same aggregation applies to all later experiments as well.)

Even though bLSM's merge scheduler prevents writes from being stalled, the instantaneous write throughput still exhibits a large variance with regular temporary peaks. Recall that bLSM uses the merge progress at each level to control its in-memory write speed. After the component C_1 is full and becomes C_1', the original C_1 will be empty and will have much shorter merge times. This will temporarily increase the in-memory write speed of bLSM, which then quickly drops as C_1 grows larger and larger. Moreover, the Zipf update workload increases the write throughput only because updated entries can be reclaimed earlier, but the overall variance performance trends are still the same.

Running Phase. Based on the maximum write throughput measured in the testing phase, we then used a constant data arrival process (95% of the maximum) in the running phase to evaluate bLSM's behavior. Figure 6b shows the instantaneous write throughput of bLSM under the uniform and Zipf update workloads. bLSM maintains a sustained write throughput during the initial period of the experiment, but later has to slow down its in-memory write rate periodically due to background merge pressure. Figure 6c further shows the resulting percentile write and processing latencies. The processing latency measures only the time for the LSM-tree to process a write, while the write latency includes both the write's queuing time and processing time. By slowing down the in-memory write rate, bLSM indeed bounds the processing latency. However, the write latency is much larger because writes must be queued when they cannot be processed immediately. This suggests that simply bounding the maximum processing latency is far from sufficient; it is important to minimize the variance in an LSM-tree's processing rate to minimize write latencies.

5. FULL MERGES

In this section, we explore the scheduling choices of LSM-trees with full merges and then evaluate the impact of merge scheduling on write stalls using our two-phase approach.

5.1 Merge Scheduling for Full Merges

We first introduce some useful notation for use throughout our analysis in Table 1. To simplify the analysis, we will ignore the I/O cost of flushes since merges consume most of the I/O bandwidth.

5.1.1 Component Constraint

To provide acceptable query performance and space utilization, the total number of disk components of an LSM-tree must be bounded. We call this upper bound the *component constraint*, and it can be enforced either *locally* or *globally*. A local constraint specifies the maximum number of disk components per level. For example, bLSM [50] uses a local constraint to allow at most two components per

Figure 6: Two-Phase Evaluation of bLSM

Table 1: List of notation used in this paper

Term	Definition	Unit
T	size ratio of the merge policy	
L	the number of levels in an LSM-tree	
M	memory component size	entries
B	I/O bandwidth	entries/s
μ	write arrival rate	entries/s
W	write throughput of an LSM-tree	entries/s

level. A global constraint instead specifies the maximum number of disk components across all levels. Here we argue that global component constraints will better minimize write stalls. In addition to external factors, such as deletes or shifts in write patterns, the merge time at each level inherently varies for leveling since the size of the component at Level i varies from 0 to $(T-1) \cdot M \cdot T^{i-1}$. Because of this, bLSM cannot provide a high yet stable write throughput over time. Global component constraints will better absorb this variance and minimize the write stalls .

It remains a question how to determine the maximum number of disk components for the component constraint. In general, tolerating more disk components will increase the LSM-tree's ability to reduce write stalls and absorb write bursts, but it will decrease query performance and space utilization. Given the negative impact of stalls on write latencies, one solution is to tolerate a sufficient number of disk components to avoid write stalls while the worst-case query performance and space utilization are still bounded. For example, one conservative constraint would be to tolerate twice the expected number of disk components, e.g., $2 \cdot L$ components for leveling and $2 \cdot T \cdot L$ components for tiering.

5.1.2 Interaction with Writes

When the component constraint is violated, the processing of writes by an LSM-tree has to be slowed down or stopped. Existing LSM-tree implementations [5, 7, 50] prefer to gracefully slow down the in-memory write rate by adding delays to some writes. This approach reduces the maximum processing latency, as large pauses are broken down into many smaller ones, but the overall processing rate of an LSM-tree, which depends on the I/O cost of each write, is not affected. Moreover, this approach will result in an even larger queuing latency. There may be additional considerations for gracefully slowing down writes, but we argue that processing writes as quickly as possible minimizes the overall write latency, as stated by the following theorem. See [41] for detailed proofs for all theorems.

THEOREM 1. *Given any data arrival process and any LSM-tree, processing writes as quickly as possible minimizes the latency of each write.*

PROOF SKETCH. Consider two merge schedulers S and S' which only differ in that S may add arbitrary delays to writes while S' processes writes as quickly as possible. For each write request r, r must be completed by S' no later than S because the LSM-tree has the same processing rate but S adds some delays to writes.

It should be noted that Theorem 1 only considers write latencies. By processing writes as quickly as possible, disk components can stack up more quickly (up to the component constraint), which may negatively impact query performance. Thus, a better approach may be to increase the write processing rate, e.g., by changing the structure of the LSM-tree. We leave the exploration of this direction as future work.

5.1.3 Degree of Concurrency

A merge policy can often create multiple merge operations simultaneously. For full merges, we can show that a single-threaded scheduler that executes one merge at a time is not sufficient for minimizing write stalls. Consider a merge operation at Level i. For leveling, the merge time varies from 0 to $\frac{M \cdot T^i}{B}$ because the size of the component at Level i varies from 0 to $(T-1) \cdot M \cdot T^{i-1}$. For tiering, each component has size $M \cdot T^{i-1}$ and merging T components thus takes time $\frac{M \cdot T^i}{B}$. Suppose the arrival rate is μ. Without concurrent merges, there would be $\frac{\mu}{M} \cdot \frac{M \cdot T^i}{B} = \frac{\mu \cdot T^i}{B}$ newly flushed components added while this merge operation is being executed, assuming that flushes can still proceed.

Our two-phase evaluation approach chooses the maximum write throughput of an LSM-tree as the arrival rate μ. For leveling, the maximum write throughput is approximately $W_{level} = \frac{2 \cdot B}{T \cdot L}$, as each entry is merged $\frac{T}{2}$ times per level. For tiering, the maximum write throughput is approximately $W_{tier} = \frac{B}{L}$, as each entry is merged only once per level. By substituting W_{level} and W_{tier} for μ, one needs to tolerate at least $\frac{2 \cdot T^{i-1}}{L}$ flushed components for leveling and $\frac{T^i}{L}$ flushed components for tiering to avoid write stalls. Since the term T^i grows exponentially, a large number of flushed components will have to be tolerated when a large disk component is being merged. Consider the leveling merge policy with a size ratio of 10. To merge a disk component at Level 5, approximately $\frac{2 \cdot 10^4}{5} = 4000$ flushed components would need to be tolerated, which is highly unacceptable.

Clearly, concurrent merges must be performed to minimize write stalls. When a large merge is being processed,

smaller merges can still be completed to reduce the number of components. By the definition of the tiering and leveling merge policies, there can be at most one active merge operation per level. Thus, given an LSM-tree with L levels, at most L merge operations can be scheduled concurrently.

5.1.4 I/O Bandwidth Allocation

Given multiple active merge operations, the merge scheduler must further decide how to allocate I/O bandwidth to these operations. A heuristic used by existing systems [2, 4, 7] is to allocate I/O bandwidth fairly (evenly) to all ongoing merges. We call this the *fair* scheduler. The fair scheduler ensures that all merges at different levels can proceed, thus eliminating potential starvation. Recall that write stalls occur when an LSM-tree has too many disk components, thus violating the component constraint. It is unclear whether or not the fair scheduler can minimize write stalls by minimizing the number of disk components over time.

Recall that both the leveling and tiering merge policies always merge the same number of disk components at once. We propose a novel *greedy* scheduler that always allocates the full I/O bandwidth to the merge operation with the smallest remaining number of bytes. The greedy scheduler has a useful property that it minimizes the number of disk components over time for a given set of merge operations.

THEOREM 2. *Given any set of merge operations that process the same number of disk components and any I/O bandwidth budget, the greedy scheduler minimizes the number of disk components at any time instant.*

PROOF SKETCH. Consider an arbitrary scheduler S and the greedy scheduler S'. Given N merge operations, we can show that S' always completes the i-th ($1 \le i \le N$) merge operation no later than S. This can be done by noting that S' always processes the smallest merge operation first.

Theorem 2 only considers a set of statically created merge operations. This conclusion may not hold in general because sometimes completing a large merge may enable the merge policy to create smaller merges, which can then reduce the number of disk components more quickly. Because of this, there actually exists no merge scheduler that can always minimize the number of disk components over time, as stated by the following theorem. However, as we will see in our later evaluation, the greedy scheduler is still a very effective heuristic for minimizing write stalls.

THEOREM 3. *Given any I/O bandwidth budget, no merge scheduler can minimize the number of disk components at any time instant for any data arrival process and any LSM-tree for a deterministic merge policy where all merge operations process the same number of disk components.*

PROOF SKETCH. Consider an LSM-tree that has created a small merge M_S and a large merge M_L. Completing M_L allows the LSM-tree to create a new small merge M'_S that is smaller than M_S. Consider two merge schedulers S_1 and S_2, where S_1 first processes M_S and then M_L, and S_2 first processes M_L and then M'_S. It can be shown that S_1 has the earliest completion time for the first merge and S_2 has the earliest completion time for the second merge, but no merge scheduler can outperform both S_1 and S_2.

5.1.5 Putting Everything Together

Based on the discussion of each scheduling choice, we now summarize the proposed greedy scheduler. The greedy scheduler enforces a global component constraint with a suf-ficient number of disk components, e.g., twice the expected number of components of an LSM-tree, to minimize write stalls while ensuring the stability of the LSM-tree. It processes writes as quickly as possible and only stops the processing of writes when the component constraint is violated. The greedy scheduler performs concurrent merges but allocates the full I/O bandwidth to the merge operation with the smallest remaining bytes. Whenever a merge operation is created or completed, the greedy scheduler is notified to find the smallest merge to execute next. Thus, a large merge can be interrupted by a newly created smaller merge. However, in general one cannot exactly know which merge operation requires the least amount of I/O bandwidth until the new component is fully produced. To handle this, the smallest merge operation can be approximated by using the number of remaining pages of the merging components.

Under the greedy scheduler, larger merges may be starved at times since they receive lower priority. This has a few implications. First, during normal user workloads, such starvation can only occur if the arrival rate is temporarily faster than the processing rate of an LSM-tree. Given the negative impact of write stalls on write latencies, it can actually be beneficial to temporarily delay large merges so that the system can better absorb write bursts. Second, the greedy scheduler should not be used in the testing phase because it would report a higher but unsustainable write throughput due to such starved large merges.

Finally, our discussions of the greedy scheduler as well as the single-threaded scheduler are based on an important assumption that a single merge operation is able to fully utilize the available I/O bandwidth budget. Otherwise, multiple merges must be executed at the same time. It is straightforward to extend the greedy scheduler to execute the smallest k merge operations, where k is the degree of concurrency needed to fully utilize the I/O bandwidth budget.

5.2 Experimental Evaluation

We now experimentally evaluate the write stalls of LSM-trees using our two-phase approach. We discuss the specific experimental setup followed by the detailed evaluation, including the impact of merge schedulers on write stalls, the benefit of enforcing the component constraint globally and of processing writes as quickly as possible, and the impact of merge scheduling on query performance.

5.2.1 Experimental Setup

All experiments in this section were performed using AsterixDB with the general setup described in Section 3. Unless otherwise noted, the size ratio of leveling was set at 10, which is a commonly used configuration in practice [5, 7]. For the experimental dataset with 100 million unique records, this results in a three-level LSM-tree, where the last level is nearly full. For tiering, the size ratio was set at 3, which leads to better write performance than leveling without sacrificing too much on query performance. This ratio results in an eight-level LSM-tree.

We evaluated the single-threaded scheduler (Section 5.1.3), the fair scheduler (Section 5.1.3), and the proposed greedy scheduler (Section 5.1.5). The single-threaded scheduler only executes one merge at a time using a single thread. Both the fair and greedy schedulers are concurrent schedulers that execute each merge using a separate thread. The difference is that the fair scheduler allocates the I/O bandwidth to all

Figure 7: Testing Phase: Instantaneous Write Throughput

ongoing merges evenly, while the greedy scheduler always allocates the full I/O bandwidth to the smallest merge. To minimize flush stalls, a flush operation is always executed in a separate thread and receives higher I/O priority. Unless otherwise noted, all three schedulers enforce global component constraints and process writes as quickly as possible. The maximum number of disk components is set at twice the expected number of disk components for each merge policy. Each experiment was performed under both the uniform and Zipf update workloads. Since the Zipf update workload had little impact on the overall performance trends, except that it led to higher write throughput, its experiment results are omitted here for brevity.

5.2.2 Testing Phase

During the testing phase, we measured the maximum write throughput of an LSM-tree by writing as much data as possible. In general, alternative merge schedulers have little impact on the maximum write throughput since the I/O bandwidth budget is fixed, but their measured write throughput may be different due to the finite experimental period.

Figures 7a and 7b shows the instantaneous write throughout of LSM-trees using different merge schedulers for tiering and leveling. Under both merge policies, the single-threaded scheduler regularly exhibits long pauses, making its write throughput vary over time. The fair scheduler exhibits a relatively stable write throughput over time since all merge levels can proceed at the same rate. With leveling, its write throughput still varies slightly over time since the component size at each level varies. The greedy scheduler appears to achieve a higher write throughput than the fair scheduler by starving large merges. However, this higher write throughput eventually drops when no small merges can be scheduled. For example, the write throughput with tiering drops slightly at 1100s and 4000s, and there is a long pause from 6000s to 7000s with leveling. This result confirms that the fair scheduler is more suitable for testing the maximum write throughput of an LSM-tree, as merges at all levels can proceed at the same rate. In contrast, the single-threaded

scheduler incurs many long pauses, causing a large variance in the measured write throughput. The greedy scheduler provides a higher write throughput by starving large merges, which would be undesirable at runtime.

5.2.3 Running Phase

Turning to the running phase, we used a constant data arrival process, configured based on 95% of the maximum write throughput measured by the fair scheduler, to evaluate the write stalls of LSM-trees.

LSM-trees can provide a stable write throughput. We first evaluated whether LSM-trees with different merge schedulers can support a high write throughput with low write latencies. For each experiment, we measured the instantaneous write throughput and the number of disk components over time as well as percentile write latencies.

The results for tiering are shown in Figure 8. Both the fair and greedy schedulers are able to provide stable write throughputs and the total number of disk components never reaches the configured threshold. The greedy scheduler also minimizes the number of disk components over time. The single-threaded scheduler, however, causes a large number of write stalls due to the blocking of large merges, which confirms our previous analysis. Because of this, the single-threaded scheduler incurs large percentile write latencies. In contrast, both the fair and greedy schedulers provide small write latencies because of their stable write throughput. Figure 9 shows the corresponding results for leveling. The single-threaded scheduler again performs poorly, causing a lot of stalls and thus large write latencies. Due to the inherent variance of merge times, the fair scheduler alone cannot provide a stable write throughput; this results in relatively large write latencies. In contrast, the greedy scheduler avoids write stalls by always minimizing the number of components, which results in small write latencies.

This experiment confirms that LSM-trees can achieve a stable write throughput with a relatively small performance variance. Moreover, the write stalls of an LSM-tree heavily depend on the design of the merge scheduler.

Impact of Size Ratio. To verify our findings on LSM-trees with different shapes, we further carried out a set of experiments by varying the size ratio from 2 to 10 for both tiering and leveling. For leveling, we applied the dynamic level size optimization [30] so that the largest level remains almost full by slightly modifying the size ratio between Levels 0 and 1. This optimization maximizes space utilization without impacting write or query performance.

During the testing phase, we measured the maximum write throughput for each LSM-tree configuration using the fair scheduler, which is shown in Figure 10a. In general, a larger size ratio increases write throughput for tiering but decreases write throughput for leveling because it decreases the merge frequency of tiering but increases that of leveling. During the running phase, we evaluated the 99% percentile write latency for each LSM-tree configuration using constant data arrivals, which is shown in Figure 10b. With tiering, both the fair and greedy schedulers are able to provide a stable write throughput with small write latencies. With leveling, the fair scheduler causes large write latencies when the size ratio becomes larger, as we have seen before. In contrast, the greedy scheduler is always able to provide a stable write throughput along with small write latencies. This again confirms that LSM-trees, despite their size ratios,

**Figure 8: Running Phase of Tiering Merge Policy (95% Load)**

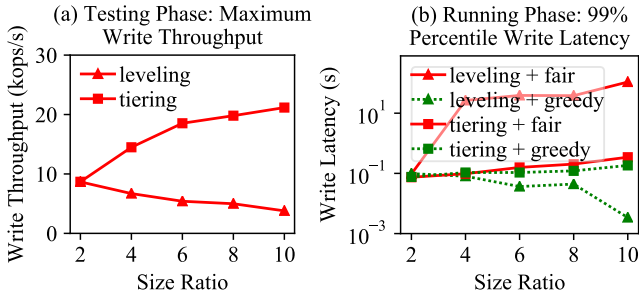**Figure 9: Running Phase of Leveling Merge Policy (95% Load)**

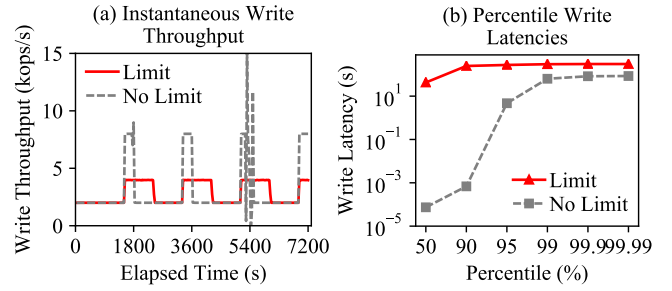**Figure 10: Impact of Size Ratio on Write Stalls**

**Figure 12: Running Phase with Burst Data Arrivals**

Figure 11: Impact of Enforcing Component Constraints on Percentile Write Latencies

can provide a high write throughput with a small variance with an appropriately chosen merge scheduler.

Benefit of Global Component Constraints. We next evaluated the benefit of global component constraints in terms of minimizing write stalls. We additionally included a variation of the fair and greedy schedulers that enforces local component constraints, that is, 2 components per level for leveling and $2 \cdot T$ components per level for tiering.

The resulting write latencies are shown in Figure 11. In general, local component constraints have little impact on

tiering since its merge time per level is relatively stable. However, the resulting write latencies for leveling become much large due to the inherent variance of its merge times. Moreover, local component constraints have a larger negative impact on the greedy scheduler. The greedy scheduler prefers small merges, which may not be able to complete due to possible violations of the constraint at the next level. This in turn causes longer stalls and thus larger percentile write latencies. In contrast, global component constraints better absorb these variances, reducing the write latencies.

Benefits of Processing Writes As Quickly As Possible. We further evaluated the benefit of processing writes as quickly as possible. We used the leveling merge policy with a bursty data arrival process that alternates between a normal arrival rate of 2000 records/s for 25 minutes and a high arrival rate of 8000 records/s for 5 minutes. We evaluated two variations of the greedy scheduler. The first variation processes writes as quickly as possible (denoted as "No Limit"), as we did before. The second variation enforces a maximum in-memory write rate of 4000 records/s (denoted as "Limit") to avoid write stalls.

The instantaneous write throughput and the percentile write latencies of the two variations are shown in Figures

**Figure 13: Instantaneous Query Throughput of Tiering Merge Policy**

Figure 14: Instantaneous Query Throughput of Leveling Merge Policy

12a and 12b respectively. As Figure 12a shows, delaying writes avoids write stalls and the resulting write throughput is more stable over time. However, this causes larger write latencies (Figure 12b) since delayed writes must be queued. In contrary, writing as quickly as possible causes occasional write stalls but still minimizes overall write latencies. This confirms our previous analysis that processing writes as quickly as possible minimizes write latencies.

Impact on Query Performance. Finally, since the point of having data is to query it, we evaluated the impact of the fair and greedy schedulers on concurrent query performance. We evaluated three types of queries, namely point lookups, short scans, and long scans. A point lookup accesses 1 record given a primary key. A short scan query accesses 100 records and a long scan query accesses 1 million records. In each experiment, we executed one type of query concurrently with concurrent updates with constant arrival rates as before. To maximize query performance while ensuring that LSM flush and merge operations receive enough I/O bandwidth, we used 8 query threads for point lookups and short scans and used 4 query threads for long scans.

The instantaneous query throughput under tiering and leveling is depicted in Figures 13 and 14 respectively. Due to space limitations, we omit the average query latency, which can be computed by dividing the number of query threads by the query throughput. As the results show, leveling has similar point lookup throughput to tiering because Bloom filters are able to filter out most unnecessary I/Os, but it has much better range query throughput than tiering. The greedy scheduler always improves query performance by minimizing the number of components. Among the three types of queries, point lookups and short scans benefit more from the greedy scheduler since these two types of queries are more sensitive to the number of disk components. In contrast, long scans incur most of their I/O cost at the largest level. Moreover, tiering benefits more from the greedy sched-

uler than leveling because tiering has more disk components. Note that with leveling, there is a drop in query throughput under the fair scheduler at around 5400s, even though there is little difference in the number of disk components between the fair and greedy schedulers. This drop is caused by write stalls during that period, as was seen in the instantaneous write throughput of Figure 9a. After the LSM-tree recovers from write stalls, it attempts to write as much data as possible to catch up, which negative impacts query performance.

We additionally evaluated the impact of forcing SSD writes regularly on query performance [41]. Even though this optimization has some small negative impact on query throughput, it significantly reduces the percentile latencies of small queries, e.g., point lookups and small scans, by 5x to 10x because the large forces at the end of merges are instead broken down into many smaller ones.

6. PARTITIONED MERGES

We now examine the write stall behavior of partitioned LSM-trees using our two-phase approach. In a partitioned LSM-tree, a large disk component is range-partitioned into multiple small files and each merge operation only processes a small number of files with overlapping ranges. Since merges always happen immediately once a level is full, a single-threaded scheduler could be sufficient to minimize write stalls. In the reminder of this section, we will evaluate LevelDB's single-threaded scheduler.

6.1 LevelDB's Merge Scheduler

LevelDB's merge scheduler is single-threaded. It computes a score for each level and selects the level with the largest score to merge. Specifically, the score for Level 0 is computed as the total number of flushed components divided by the minimum number of flushed components to merge. For a partitioned level (1 and above), its score is defined as the total size of all files at this level divided by the
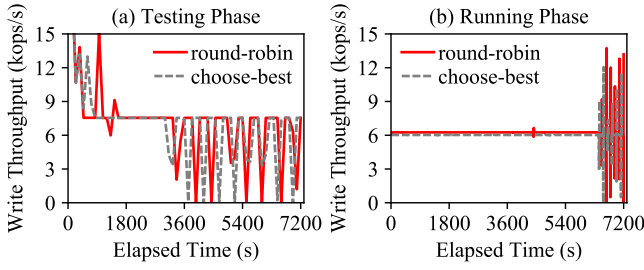
Figure 15: Instantaneous Write Throughput under Two-Phase Evaluation of Partitioned LSM-tree

Figure 16: Problem of Score-Based Merge Scheduler

configured maximum size. A merge operation is scheduled if the largest score is at least 1, which means that the selected level is full. If a partitioned level is chosen to merge, LevelDB selects the next file to merge in a round-robin way.

LevelDB only restricts the number of flushed components at Level 0. By default, the minimum number of flushed components to merge is 4. The processing of writes will be slowed down or stopped of the number of flushed component reaches 8 and 12 respectively. Since we have already shown in Section 5.1.2 that processing writes as quickly as possible reduces write latencies, we will only use the stop threshold (12) in our evaluation.

Experimental Evaluation. We have implemented LevelDB's partitioned leveling merge policy and its merge scheduler inside AsterixDB for evaluation. Similar to LevelDB, the minimum number of flushed components to merge was set at 4 and the stop threshold was set at 12 components. Unless otherwise noted, the maximum size of each file was set at 64MB. The memory component size was set at 128MB and the base size of Level 1 was set at 1280MB. The size ratio was set at 10. For the experimental dataset with 100 million records, this results in a 4-level LSM-tree where the largest level is nearly full. To minimize write stalls caused by flushes, we used two memory components and a separate flush thread. We further evaluated the impact of two widely used merge selection strategies on write stalls. The round-robin strategy chooses the next file to merge in a round-robin way. The choose-best strategy [52] chooses the file with the fewest overlapping files at the next level.

We used our two-phase approach to evaluate this partitioned LSM-tree design. The instantaneous write throughput during the testing phase is shown in Figure 15a, where the write throughput of both strategies decreases over time due to more frequent stalls. Moreover, under the uniform update workload, the alternative selection strategies have little impact on the overall write throughput, as reported in [37]. During the testing phase, we used a constant arrival process to evaluate write stalls. The instantaneous write throughput of both strategies is shown in Figure 15b. As the result shows, in both cases write stalls start to occur after time 6000s. This suggests that the measured write throughput during the testing phase is not sustainable.

6.2 Measuring Sustainable Write Throughput

One problem with LevelDB's score-based merge scheduler is that it merges as many components at Level 0 as possible at once. To see this, suppose that the minimum number of mergeable components at Level L_0 is T_0 and that the maximum number of components at Level 0 is T_0'. During the testing phase, where writes pile up as quickly as possible, the merge scheduler tends to merge the maximum possible number of components T_0' instead of just T_0 at once. Because of this, the LSM-tree will eventually transit from the expected shape (Figure 16a) to the actual shape (Figure 16b), where T is the size ratio of the partitioned levels. Note that the largest level is not affected since its size is determined by the number of unique entries, which is relatively stable. Even though this elastic design dynamically increases the processing rate as needed, it has the following problems.

Unsustainable Write Throughput. The measured maximum write throughput is achieved by merging T_0' flushed components at Level 0 at once. However, this is likely to cause write stalls during the running phase since flushes cannot further proceed.

Suboptimal Trade-Offs. The LSM-tree structure in Figure 16b is no longer making optimal performance trade-offs since the size ratios between its adjacent levels are not the same anymore [44]. By adjusting the sizes of intermediate levels so that adjacent levels have the same size ratio, one can improve both write throughput and space utilization without affecting query performance.

Low Space Utilization. One motivation for industrial systems to adopt partitioned LSM-trees is their higher space utilization [30]. However, the LSM-tree depicted in Figure 16b violates this performance guarantee because the ratio of wasted space increases from $1/T$ to $T_0'/T_0 \cdot 1/T$.

Because of these problems, the measured maximum write throughput cannot be used in the long-term. We propose a simple solution to address these problems. During the testing phase, we always merge exactly T_0 components at Level 0. This ensures that merge preferences will be given equally to all levels so that the LSM-tree will stay in the expected shape (Figure 16a). Then, during the running phase, the LSM-tree can elastically merge more components at Level 0 as needed to absorb write bursts.

To verify the effectiveness of the proposed solution, we repeated the previous experiments on the partitioned LSM-tree. During the testing phase, the LSM-tree always merged 4 components at Level 0 at once. The measured instantaneous write throughput is shown in Figure 17a, which is 30% lower than that of the previous experiment. During the running phase, we used a constant arrival process based on this lower write throughput. The resulting instantaneous write throughput is shown in Figure 17b, where the LSM-tree successfully maintains a sustainable write throughput without any write stalls, which in turn results in low write latencies (not shown in the figure). This confirms that LevelDB's single-threaded scheduler is sufficient to minimize write stalls, given that a single merge thread can fully utilize the I/O bandwidth budget.

After fixing the unsustainable write throughput problem of LevelDB, we further evaluated the impact of partition size on the write stalls of partitioned LSM-trees. In this

Figure 17: Instantaneous Write Throughput under Two-Phase Evaluation of Partitioned LSM-tree with the Proposed Solution


Figure 18: Impact of Partition Size on Write Stalls

experiment, we varied the size of each partitioned file from 8MB to 32GB so that partitioned merges effectively transit into full merges. The maximum write throughput during the running phase and the 99th percentile write latencies during the testing phase are shown in Figures 18a and 18b respectively. Even though the partition size has little impact on the overall write throughput, a large partition size can cause large write latencies since we have shown in Section 5 that a single-threaded scheduler is insufficient to minimize write stalls for full merges. Most implementations of partitioned LSM-trees today already choose a small partition size to bound the temporary space occupied by merges. We see here that one more reason to do so is to minimize write stalls under a single-threaded scheduler.

7. LESSONS AND INSIGHTS

Having studied and evaluated the write stall problem for various LSM-tree designs, here we summarize the lessons and insights observed from our evaluation.

The LSM-tree's write latency must be measured properly. The out-of-place update nature of LSM-trees has introduced the write stall problem. Throughout our evaluation, we have seen cases where one can obtain a higher but unsustainable write throughput. For example, the greedy scheduler would report a higher write throughput by starving large merges, and LevelDB's merge scheduler would report a higher but unsustainable write throughput by dynamically adjusting the shape of the LSM-tree. Based on our findings, we argue that in addition to the testing phase, used by existing LSM research, an extra running phase must be performed to evaluate the usability of the measured maximum write throughput. Moreover, the write latency must be measured properly due to queuing. One solution is to use the proposed two-phase evaluation approach to evaluate

the resulting write latencies under high utilization, where the arrival rate is close to the processing rate.

Merge scheduling is critical to minimizing write stalls. Throughout our evaluation of various LSM-tree designs, including bLSM [50], full merges, and partitioned merges, we have seen that merge scheduling has a critical impact on write stalls. Comparing these LSM-tree designs in general depends on many factors and is beyond the scope of this paper; here we have focused on how to minimize write stalls for each LSM-tree design.

bLSM [50], an instance of full merges, introduces a sophisticated spring-and-gear merge scheduler to bound the processing latency of LSM-trees. However, we found that bLSM still has large variances in its processing rate, leading to large write latencies under high arrival rates. Among the three evaluated schedulers, namely single-threaded, fair, and greedy, the single-threaded scheduler should not be used in practical systems due to the long stalls caused by large merges. The fair scheduler should be used when measuring the maximum throughput because it provides fairness to all merges. The greedy scheduler should be used at runtime since it better minimizes the number of disk components, both reducing write stalls and improving query performance. Moreover, as an important design choice, global component constraints better minimizes write stalls.

Partitioned merges simplify merge scheduling by breaking large merges into many smaller ones. However, we found a new problem that the measured maximum write throughput of LevelDB is unsustainable because it dynamically adjusts the size ratios under write-intensive workloads. After fixing this problem, a single-threaded scheduler with a small partition size, as used by LevelDB, is sufficient for delivering low write latencies under high utilization. However, fixing this problem reduced the maximum write throughput of LevelDB roughly one-third in our evaluation.

For both full and partitioned merges, processing writes as quickly as possible better minimizes write latencies. Finally, with proper merge scheduling, all LSM-tree designs can indeed minimize write stalls by delivering low write latencies under high utilizations.

8. CONCLUSION

In this paper, we have studied and evaluated the write stall problem for various LSM-tree designs. We first proposed a two-phase approach to use in evaluating the impact of write stalls on percentile write latencies using a combination of closed and open system testing models. We then identified and explored the design choices for LSM merge schedulers. For full merges, we proposed a greedy scheduler that minimizes write stalls. For partitioned merges, we found that a single-threaded scheduler is sufficient to provide a stable write throughput but that the maximum write throughput must be measured properly. Based on these findings, we have shown that performance variance must be considered together with write throughput to ensure the actual usability of the measured throughput.

9. REFERENCES

[1] AsterixDB. https://asterixdb.apache.org/.

[2] Cassandra. http://cassandra.apache.org/.

[3] Compaction stalls: something to make better in RocksDB. http://smalldatum.blogspot.com/2017/01/compaction-stalls-something-to-make.html.

[4] HBase. https://hbase.apache.org/.

[5] LevelDB. http://leveldb.org/.

[6] Read- and latency-optimized log structured merge tree. https://github.com/sears/bLSM/.

[7] RocksDB. http://rocksdb.org/.

[8] SyllaDB. https://www.scylladb.com/.

[9] Tarantool. https://www.tarantool.io/.

[10] TPC-C. http://www.tpc.org/tpcc/.

[11] WiredTiger. http://www.wiredtiger.com/.

[12] YCSB change log. https://github.com/brianfrankcooper/YCSB/blob/master/core/CHANGES.md.

[13] M. Y. Ahmad and B. Kemme. Compaction management in distributed key-value datastores. *PVLDB*, 8(8):850–861, 2015.

[14] S. Alsubaiee et al. AsterixDB: A scalable, open source BDMS. *PVLDB*, 7(14):1905–1916, 2014.

[15] M. Armbrust et al. Piql: Success-tolerant query processing in the cloud. *PVLDB*, 5(3):181–192, 2011.

[16] M. Armbrust et al. Generalized scale independence through incremental precomputation. In *ACM SIGMOD*, pages 625–636, 2013.

[17] O. Balmau et al. FloDB: Unlocking memory in persistent key-value stores. In *European Conference on Computer Systems (EuroSys)*, pages 80–94, 2017.

[18] O. Balmau et al. TRIAD: Creating synergies between memory, disk and log in log structured key-value stores. In *USENIX Annual Technical Conference (ATC)*, pages 363–375, 2017.

[19] B. H. Bloom. Space/time trade-offs in hash coding with allowable errors. *CACM*, 13(7):422–426, July 1970.

[20] G. Candea et al. A scalable, predictable join operator for highly concurrent data warehouses. *PVLDB*, 2(1):277–288, 2009.

[21] Z. Cao et al. On the performance variation in modern storage stacks. In *USENIX Conference on File and Storage Technologies (FAST)*, pages 329–343, 2017.

[22] M. J. Carey. AsterixDB mid-flight: A case study in building systems in academia. In *ICDE*, pages 1–12, 2019.

[23] S. Chaudhuri et al. Variance aware optimization of parameterized queries. In *ACM SIGMOD*, pages 531–542. ACM, 2010.

[24] B. F. Cooper et al. Benchmarking cloud serving systems with YCSB. In *ACM SoCC*, pages 143–154, 2010.

[25] N. Dayan et al. Monkey: Optimal navigable key-value store. In *ACM SIGMOD*, pages 79–94, 2017.

[26] N. Dayan et al. Optimal Bloom filters and adaptive merging for LSM-trees. *ACM TODS*, 43(4):16:1–16:48, Dec. 2018.

[27] N. Dayan and S. Idreos. Dostoevsky: Better space-time trade-offs for LSM-tree based key-value stores via adaptive removal of superfluous merging. In *ACM SIGMOD*, pages 505–520, 2018.

[28] N. Dayan and S. Idreos. The log-structured merge-bush & the wacky continuum. In *ACM SIGMOD*, pages 449–466, 2019.

[29] J. Dean and L. A. Barroso. The tail at scale. *CACM*, 56:74–80, 2013.

[30] S. Dong et al. Optimizing space amplification in RocksDB. In *CIDR*, volume 3, page 3, 2017.

[31] R. Grover and M. J. Carey. Data ingestion in AsterixDB. In *EDBT*, pages 605–616, 2015.

[32] M. Harchol-Balter. *Performance modeling and design of computer systems: queueing theory in action.* Cambridge University Press, 2013.

[33] G. Huang et al. X-Engine: An optimized storage engine for large-scale E-commerce transaction processing. In *ACM SIGMOD*, pages 651–665, 2019.

[34] J. Huang et al. Statistical analysis of latency through semantic profiling. In *European Conference on Computer Systems (EuroSys)*, pages 64–79, 2017.

[35] J. Huang et al. A top-down approach to achieving performance predictability in database systems. In *ACM SIGMOD*, pages 745–758, 2017.

[36] Y. Li et al. Tree indexing on solid state drives. *PVLDB*, 3(1-2):1195–1206, 2010.

[37] H. Lim et al. Towards accurate and fast evaluation of multi-stage log-structured designs. In *USENIX Conference on File and Storage Technologies (FAST)*, pages 149–166, 2016.

[38] L. Lu et al. WiscKey: Separating keys from values in SSD-conscious storage. In *USENIX Conference on File and Storage Technologies (FAST)*, pages 133–148, 2016.

[39] C. Luo and M. J. Carey. Efficient data ingestion and query processing for LSM-based storage systems. *PVLDB*, 12(5):531–543, 2019.

[40] C. Luo and M. J. Carey. LSM-based storage techniques: a survey. *The VLDB Journal*, 2019.

[41] C. Luo and M. J. Carey. On performance stability in LSM-based storage systems. *CoRR*, abs/1906.09667, 2019.

[42] C. Luo et al. Umzi: Unified multi-zone indexing for large-scale HTAP. In *EDBT*, pages 1–12, 2019.

[43] F. Mei et al. SifrDB: A unified solution for write-optimized key-value stores in large datacenter. In *ACM SoCC*, pages 477–489, 2018.

[44] P. O'Neil et al. The log-structured merge-tree (LSM-tree). *Acta Inf.*, 33(4):351–385, 1996.

[45] M. A. Qader et al. A comparative study of secondary indexing techniques in LSM-based NoSQL databases. In *ACM SIGMOD*, pages 551–566, 2018.

[46] P. Raju et al. PebblesDB: Building key-value stores using fragmented log-structured merge trees. In *ACM SOSP*, pages 497–514, 2017.

[47] V. Raman et al. Constant-time query processing. In *ICDE*, pages 60–69, 2008.

[48] K. Ren et al. SlimDB: A space-efficient key-value storage engine for semi-sorted data. *PVLDB*, 10(13):2037–2048, 2017.

[49] R. Sears and E. Brewer. Stasis: Flexible transactional storage. In *Symposium on Operating Systems Design and Implementation (OSDI)*, pages 29–44, 2006.

[50] R. Sears and R. Ramakrishnan. bLSM: A general purpose log structured merge tree. In *ACM SIGMOD*, pages 217–228, 2012.

[51] D. Teng et al. LSbM-tree: Re-enabling buffer caching in data management for mixed reads and writes. In *IEEE International Conference on Distributed Computing Systems (ICDCS)*, pages 68–79, 2017.

[52] R. Thonangi and J. Yang. On log-structured merge for solid-state drives. In *ICDE*, pages 683–694, 2017.

[53] P. Unterbrunner et al. Predictable performance for unpredictable workloads. *PVLDB*, 2(1):706–717, 2009.

[54] X. Wang and M. J. Carey. An IDEA: An ingestion framework for data enrichment in AsterixDB. *PVLDB.*, 12(11):1485–1498, 2019.

[55] T. Yao et al. A light-weight compaction tree to reduce I/O amplification toward efficient key-value stores. In *International Conference on Massive Storage Systems and Technology (MSST)*, 2017.

[56] C. Yunpeng et al. LDC: a lower-level driven compaction method to optimize SSD-oriented key-value stores. In *ICDE*, pages 722–733, 2019.

[57] Y. Zhang et al. ElasticBF: Fine-grained and elastic bloom filter towards efficient read for LSM-tree-based KV stores. In *USENIX Workshop on Hot Topics in Storage and File Systems (HotStorage)*, 2018.